



SOaC Framework Architecture

Overview

The SOaC Framework is built on a modern, microservices-inspired architecture that enables scalable, maintainable, and extensible security operations automation.

Architecture Diagram







Core Components

1. API Integration Layer

Purpose: Communicate with external security devices and collect data.

Components:

- **Base Connector:** Abstract base class with common functionality
- **Device-Specific Connectors:** Implementations for each platform
- PaloAlto Connector (PAN-OS REST API)
- EntralID Connector (Microsoft Graph API)
- SIEM Connector (Splunk/Elastic REST APIs)

Key Features:

- Connection pooling and retry logic
- Rate limiting and throttling
- Credential management
- Health checking
- Mock mode for testing

2. Data Processing Layer

Purpose: Normalize, validate, and enrich security events from different sources.

Components:

- **Event Ingestion Service:** Background service for continuous event collection
- **Normalization Engine:** Convert vendor-specific formats to common schema
- **Validation Service:** Ensure data quality and completeness
- **Enrichment Service:** Add context (geolocation, threat intel, asset info)

Common Event Schema:

```
{
  "id": "uuid",
  "source": "device_name",
  "source_type": "paloalto|entraid|siem",
  "timestamp": "2025-11-14T10:00:00Z",
  "event_type": "authentication|network|process|file",
  "severity": "critical|high|medium|low|info",
  "entities": {
    "user": "john.doe@company.com",
    "host": "LAPTOP01",
    "ip": "192.168.1.100",
    "file": "malware.exe",
    "hash": "sha256:abcd1234..."
  },
  "raw_data": {...}
}
```

3. Detection Layer

Purpose: Correlate events across time and sources to detect complex attacks.

Components:

Correlation Engine

- **Multi-Phase Detection:** Track attack progression through multiple stages
- **Entity Tracking:** Follow users, hosts, IPs across events
- **Temporal Windowing:** Correlate events within time windows (5 min - 24 hours)
- **Confidence Scoring:** Calculate confidence based on matched phases

Correlation Logic:

```
# Pseudocode for correlation
def correlate_events(events, time_window, min_phases):
    entities = extract_entities(events)

    for entity in entities:
        event_chain = get_events_by_entity(entity, time_window)
        matched_phases = match_phases(event_chain, operational_model)

        if len(matched_phases) >= min_phases:
            confidence = calculate_confidence(matched_phases)
            create_incident(entity, event_chain, confidence)
```

Operational Models

Pre-built detection patterns for common attacks:

- **Ransomware:** Delivery → Execution → Encryption → Impact
- **Data Theft:** Collection → Staging → Exfiltration
- **Intrusion:** Foothold → Privilege Abuse → Lateral Movement
- **Fraud:** Compromise → Transaction → Exfiltration
- **DoS:** Flood → Degradation → Exhaustion

Model Structure:

```
{
  "name": "Ransomware Detection",
  "phases": [
    {
      "name": "Delivery",
      "sources": ["proofpoint", "email_gateway"],
      "indicators": ["attachment with .exe", "suspicious links"]
    },
    {
      "name": "Execution",
      "sources": ["falcon", "endpoint_edr"],
      "indicators": ["powershell -enc", "suspicious process tree"]
    },
    {
      "name": "Encryption",
      "sources": ["falcon", "file_monitoring"],
      "indicators": ["mass file rename", ".locked extension"]
    }
  ],
  "correlation_fields": ["user", "computer", "ip"],
  "time_window": "60 minutes",
  "min_phases": 3
}
```

4. Incident Management Layer

Purpose: Manage security incidents through their lifecycle.

Components:

- **Incident Manager:** CRUD operations for incidents
- **Workflow Engine:** State transitions and approvals
- **Assignment Engine:** Auto-assign based on severity/type
- **Enrichment Service:** Add threat intel and context

Incident Lifecycle:

```
New → Assigned → Investigating → Contained →
Resolved → Closed
```

Incident Structure:

```
{
  "id": "INC-12345678",
  "title": "Ransomware Attack Detected",
  "severity": "critical",
  "confidence": "high",
  "status": "investigating",
  "assignee": "analyst@company.com",
  "operational_model": "ransomware",
  "matched_phases": ["delivery", "execution", "encryption"],
  "entities": {
    "user": "john.doe@company.com",
    "host": "LAPTOP01",
    "ip": "192.168.1.100"
  },
  "events": [...],
  "timeline": [...],
  "playbooks_executed": [...]
}
```

5. Response Layer (SOAR)

Purpose: Automate security response actions across multiple platforms.

Components:

Playbook Engine

- **Decision Matrix:** Determine which playbooks to execute
- **Action Executor:** Execute actions across devices
- **Approval Workflow:** Manual approval for sensitive actions
- **Rollback:** Revert actions if needed

Playbook Library

Standard playbooks for common responses:

1. Endpoint Containment

- Isolate host from network
- Kill malicious processes

- Capture memory/disk for forensics
- Tag device in CMDB

2. Identity Lockdown

- Disable user account
- Revoke active sessions
- Reset MFA
- Force password change

3. Network Containment

- Block IP/domain at firewall
- Enable packet capture
- Rate limit connections
- Create quarantine VLAN

4. Cloud Mitigation

- Revoke API keys/tokens
- Lock cloud resources
- Snapshot for forensics
- Roll back configuration

Playbook Structure:

```
{
  "name": "Ransomware Containment",
  "trigger_conditions": {
    "operational_model": "ransomware",
    "confidence": "high",
    "matched_phases": [">=3"]
  },
  "steps": [
    {
      "action": "isolate_endpoint",
      "device": "falcon_edr",
      "params": {"host": "${incident.entities.host}"}
    },
    {
      "action": "disable_account",
      "device": "entraid",
      "params": {"user": "${incident.entities.user}"}
    },
    {
      "action": "block_ip",
      "device": "paloalto_ngfw",
      "params": {"ip": "${incident.entities.ip}"}
    },
    {
      "action": "create_ticket",
      "device": "servicenow",
      "params": {"incident_id": "${incident.id}"}
    }
  ],
  "approval_required": false,
  "auto_execute": true
}
```

6. Presentation Layer

Purpose: Provide user interfaces for security operations.

Components:

REST API (FastAPI)

- **OpenAPI/Swagger Documentation:** Auto-generated API docs
- **JWT Authentication:** Stateless authentication
- **Role-Based Access Control:** Admin, Analyst, Viewer roles
- **Rate Limiting:** Prevent abuse
- **Audit Logging:** Track all API calls

Key Endpoints:

- `/api/v1/auth/*` - Authentication
- `/api/v1/devices/*` - Device management
- `/api/v1/rules/*` - Detection rules
- `/api/v1/events/*` - Event management
- `/api/v1/incidents/*` - Incident management
- `/api/v1/detection/*` - Detection engine
- `/api/v1/playbooks/*` - Playbook execution

Web UI (React + TypeScript)

- **Dashboard:** Real-time metrics and status
- **Device Management:** Configure and monitor devices
- **Rule Management:** Create and manage detection rules
- **Event Browser:** Search and filter events
- **Incident Investigation:** Full incident details and timeline
- **Operational Models:** View and configure detection patterns
- **Playbook Management:** Configure and test playbooks

7. Data Layer

Purpose: Persistent storage for all framework data.

Components:

PostgreSQL Database

Primary data store with tables for:

- **devices:** Security device configurations
- **rules:** Detection rules and mappings
- **events:** Collected security events
- **incidents:** Security incidents
- **users:** User accounts and roles
- **playbooks:** Response playbook definitions
- **audit_logs:** Complete audit trail

Redis Cache

Used for:

- **Task Queue:** Celery background jobs
- **Session Storage:** User sessions

- **Rate Limiting:** API rate limit counters
- **Caching:** Frequently accessed data

Data Flow

Event Collection Flow

1. Scheduled Task (every 5 min)
↓
2. Sync Service calls Device Connector
↓
3. Device Connector fetches events via API
↓
4. Events normalized **to** common schema
↓
5. Events **validated and enriched**
↓
6. Events **saved to PostgreSQL**
↓
7. Events sent **to** Correlation Engine

Detection Flow

1. **New** event arrives
↓
2. Correlation Engine extracts entities
↓
3. Query recent events **for** same entities
↓
4. Match events against Operational Models
↓
5. Calculate confidence score
↓
6. **If** threshold met, create Incident
↓
7. Incident triggers Decision Matrix
↓
8. Appropriate Playbooks executed
↓
9. Actions performed **on** devices
↓
10. Incident updated with results

Scalability Considerations

Horizontal Scaling

- **Backend API:** Stateless design allows multiple instances
- **Event Ingestion:** Can be split by device type
- **Correlation Engine:** Can process events in parallel
- **Database:** Read replicas for queries

Performance Optimizations

- **Database Indexing:** Optimized indexes on frequently queried fields
- **Caching:** Redis for frequently accessed data

- **Background Jobs:** Celery for async processing
- **Connection Pooling:** Reuse connections to devices and database

High Availability

- **Load Balancing:** Distribute traffic across backend instances
- **Database Replication:** PostgreSQL streaming replication
- **Redundant Workers:** Multiple Celery workers
- **Health Checks:** Kubernetes liveness/readiness probes

Security Architecture

Authentication & Authorization

- **JWT Tokens:** Short-lived access tokens
- **Refresh Tokens:** Long-lived refresh tokens
- **Role-Based Access:** Admin, Analyst, Viewer roles
- **API Keys:** For service-to-service communication

Secrets Management

- **Environment Variables:** For development
- **AWS Secrets Manager:** For AWS deployments
- **Azure Key Vault:** For Azure deployments
- **Kubernetes Secrets:** For K8s deployments

Network Security

- **TLS/SSL:** All communication encrypted
- **Network Segmentation:** Isolated network zones
- **Firewall Rules:** Whitelist-based access
- **VPN:** Secure access to management interfaces

Data Security

- **Encryption at Rest:** Database encryption
- **Encryption in Transit:** TLS for all connections
- **Credential Encryption:** Device credentials encrypted in DB
- **Audit Logging:** Complete audit trail

Technology Choices

Why FastAPI?

- Modern Python web framework
- Automatic OpenAPI documentation
- Built-in data validation (Pydantic)
- Async support for performance
- Type hints for better code quality

Why React + TypeScript?

- Modern, component-based UI
- Type safety reduces bugs

- Large ecosystem of libraries
- Excellent developer experience
- Strong community support

Why PostgreSQL?

- ACID compliance
- Powerful query capabilities
- JSON support for flexible schemas
- Strong ecosystem
- Proven reliability

Why Docker/Kubernetes?

- Consistent environments
- Easy deployment
- Horizontal scaling
- Service isolation
- Industry standard

Monitoring & Observability

Application Metrics

- Request rate and latency
- Error rates
- Event ingestion rate
- Incident creation rate
- Playbook execution success rate

Infrastructure Metrics

- CPU and memory usage
- Database connections
- Queue depth
- Network I/O
- Disk usage

Logging


- Structured JSON logs
- Centralized log aggregation
- Log levels (DEBUG, INFO, WARNING, ERROR)
- Correlation IDs for request tracking

Alerting


- Application health alerts
- Database performance alerts
- Integration failures
- Security events

Deployment Architecture

Docker Compose (Development)



Kubernetes (Production)



Future Architecture Enhancements

Version 1.1

- Message queue (RabbitMQ/Kafka) for event streaming
- Elasticsearch for event search
- Grafana for advanced dashboards
- Multi-tenancy support

Version 2.0

- AI/ML for anomaly detection
- Graph database for entity relationships
- Real-time streaming analytics
- Advanced threat hunting

For implementation details, see the [Development Guide](#) (./DEVELOPMENT.md)